

Quality Control

Sergey Naumenko

2021-05-08

Contents

Overview	1
Metadata	1
Read metrics	3
Total reads	3
Mapped reads	5
Number of genes detected	7
Gene detection saturation	9
Exonic mapping rate	10
Intronic mapping rate	12
rRNA mapping rate	14
5'->3' bias	16
Counts per gene - all genes	18
Counts per gene - protein coding genes	20
Sample similarity analysis	22
Principal component analysis (PCA) - non zero genes	22
Principal component analysis (PCA) - protein coding genes	23
Principal component analysis (PCA) - protein coding genes - er	25
PCA - protein coding genes - PRE/POST - Response	28
PRE	28
PRE - ROC	30
POST	31
POST - ROC	33
R session	34

Overview

- Principal Investigator: Beth Overmoyer
- Experiment: RNAseq_analysis_of_inflammatory_breast_cancer_hbc04141
- study 6 was excluded

Metadata

```
se <- readRDS("data/bcbio-se_all_samples.rds")
se <- se[, se$study_id != 6]
```

```

mdata <- colData(se)
mdata <- subset(mdata, subset = study_id != 6)
mdata$response <- str_replace(mdata$response, "Yes", "pCR")
mdata$response <- str_replace(mdata$response, "No", "non-pCR")
colData(se) <- mdata

mdata <- metadata(se)
mdata$metrics <- mdata$metrics %>% dplyr::filter(!sample %in% c("s3370_1", "s3373_3"))
metadata(se) <- mdata

metadata <- colData(se) %>%
  as_tibble(rownames = NULL) %>%
  dplyr::select(-batch, -phenotype)

metrics <- metadata(se)$metrics %>%
  left_join(metadata, by = c("sample" = "sample"))

metrics$date_of <- as_factor(metrics$date_of)

metadata

## # A tibble: 44 x 8
##   category date_of er      response study_id treatment tumor_percentage sample
##   <chr>      <dbl> <chr>   <chr>      <dbl> <chr>      <chr>      <chr>
## 1 pre      20180228 Positi~ pCR          1 pre      30-100      s3154~
## 2 post     20180228 Positi~ pCR          1 post     30-100      s3169~
## 3 pre      20180323 Positi~ non-pCR      3 pre      1-29       s3188~
## 4 pre      20180323 Positi~ <NA>         2 pre      30-100      s3190~
## 5 post     20180323 Positi~ non-pCR      3 post     1-29       s3193~
## 6 post     20180228 Positi~ <NA>         2 post     1-29       s3194~
## 7 pre      20180228 Positi~ pCR          4 pre      30-100      s3220~
## 8 post     20180228 Positi~ pCR          4 post     30-100      s3234~
## 9 pre      20180228 Positi~ pCR          5 pre      1-29       s3291~
## 10 post    20180228 Positi~ pCR          5 post     1-29       s3292~
## # ... with 34 more rows


```

Read metrics

Total reads

```
metrics <- metrics %>%
  group_by(date_of) %>%
  mutate(position = rank(-total_reads))


metrics %>%
  ggplot(aes(x = reorder(sample, -total_reads),
              y = total_reads/1e6L,
              fill = date_of,
              group = position)) +
  geom_bar(stat = "identity") +
  coord_flip() +
  scale_y_continuous(name = "mln reads") +
  geom_text(aes(label = floor(total_reads/1e6L)), hjust = 0, nudge_y = 0.5)+
  xlab("sample") +
  ggtitle("Total reads")
```



Mapped reads

The number of mapped reads should correspond to the number of total reads.

```
metrics %>%  
  ggplot(aes(x = reorder(sample, -mapped_reads),  
             y = mapped_reads, fill = date_of)) +  
  geom_bar(stat = "identity") +  
  coord_flip() +  
  scale_y_continuous(name = "mln reads") +  
  geom_text(aes(label = floor(mapped_reads/1e6L)), hjust = 0, nudge_y = 0.5)+  
  xlab("sample") +  
  ggtitle("Mapped reads")
```




```
## Mapping rate
#The genomic mapping rate represents the percentage of reads mapping to the reference genome. Low mapping
#{r plot_mapping_rate, fig.width = 10, fig.height = 20}
#metrics %>%
#   ggplot(aes(x = reorder(sample, -mapped_reads_pct),
#               y = mapped_reads_pct, fill = tissue)) +
#       geom_bar(stat = "identity") +
#   coord_flip() +
#   geom_text(aes(label = floor(mapped_reads_pct)), hjust = 0, nudge_y = 0.5)+
#   xlab("sample") +
#   ggtitle("Mapping rate")
```

Number of genes detected

```
genes_detected <- colSums(assays(se)[["raw"]] > 0) %>% enframe()
colnames(genes_detected) <- c("sample", "n_genes")

metrics <- metrics %>%
  left_join(genes_detected, by = c("sample" = "sample"))


metrics %>%
  ggplot(aes(x = reorder(sample, -n_genes),
              y = n_genes, fill = date_of)) +
  geom_bar(stat = "identity") +
  coord_flip() +
  geom_text(aes(label = n_genes), hjust = 0, nudge_y = 0.5)+
  xlab("sample") +
  ylim(0, 50000) +
  ggtitle("Number of genes")
```



Gene detection saturation

We should observe a linear trend in the number of genes detected with the number of mapped reads, which indicates that the sample input was not overloaded.


```
metrics %>%  
  ggplot(aes(x = log10(total_reads),  
             y = n_genes,  
             color = date_of)) +  
    geom_point(size = 5) +  
    geom_text_repel(aes(label = sample), size = 10) +  
  theme(axis.text = element_text(size = 20),  
        axis.title = element_text(size = 20, face="bold"),  
        legend.text = element_text(size = rel(1))) +  
  theme_bw(base_size = 20) +  
  ggtitle("Gene saturation")
```



Exonic mapping rate

Ideally, at least 60% of total reads should map to exons.


```
metrics %>%
  ggplot(aes(x = reorder(sample, -exonic_rate),
             y = exonic_rate * 100,
             fill = date_of)) +
  geom_bar(stat = "identity") +
  geom_text(aes(label = floor(exonic_rate*100)), hjust = 0, nudge_y = 0.5) +
  xlab("sample") +
  ylab("Exonic rate %") +
  ggtitle("Exonic mapping rate") +
  coord_flip()
```



Intronic mapping rate

The majority of reads should map to exons and not introns.

```
metrics %>%  
  ggplot(aes(x = reorder(sample, -intronic_rate),  
             y = intronic_rate * 100,  
             fill = date_of)) +  
  geom_bar(stat = "identity") +  
  coord_flip() +  
  geom_text(aes(label = floor(intronic_rate*100)),  
            hjust = 0,  
            nudge_y = 0.5) +  
  xlab("sample") +  
  ylab("Intronic rate %")
```




```
ggtitle("Intronic mapping rate")
```

```
## $title  
## [1] "Intronic mapping rate"  
##  
## attr(,"class")  
## [1] "labels"
```

rRNA mapping rate

Samples should have a ribosomal RNA (rRNA) contamination rate below 10%.

```
metrics %>%  
  ggplot(aes(x = reorder(sample, -r_rna_rate),  
             y = r_rna_rate * 100,  
             fill = date_of)) +  
  geom_bar(stat = "identity") +  
  geom_text(aes(label = round(r_rna_rate*100,2)), hjust = 0, nudge_y = 0.5)+  
  xlab("sample") +  
  ylab("rRNA rate, %")+  
  ylim(0, 1) +  
  ggtitle("rRNA mapping rate") +  
  coord_flip()
```



5'→3' bias


```
metrics %>%  
  ggplot(aes(x = reorder(sample, -x5_3_bias),  
             y = x5_3_bias,  
             fill = date_of)) +  
    geom_bar(stat = "identity") +  
    geom_text(aes(label = x5_3_bias), hjust = 0, nudge_y = 0.01)+  
    xlab("sample") +  
    ylim(0, 1.5)+  
    ggtitle("5'-3' bias") +  
    coord_flip()
```




Counts per gene - all genes


We expected similar spread for every sample.

```
metadata$date_of <- as_factor(metadata$date_of)
metrics_small <- metrics %>% dplyr::select(sample, date_of)
assays(se)[["raw"]] %>%
  as_tibble() %>%
  dplyr::filter(rowSums(.)!=0) %>%
  gather(sample, counts) %>%
  left_join(metadata, by = c("sample" = "sample")) %>%
  ggplot(aes(sample, log2(counts+1), fill = date_of)) +
  geom_boxplot() +
  coord_flip() +
  ggtitle("Counts per gene, all non-zero genes")
```



Counts per gene - protein coding genes

```
protein_coding_genes <- read_csv("tables/ensembl_w_description.protein_coding.csv")
assays(se)[["raw"]] %>%
  as_tibble(rownames = "ensembl_gene_id") %>%
  dplyr::filter(ensembl_gene_id %in% protein_coding_genes$ensembl_gene_id) %>%
  dplyr::select(-ensembl_gene_id) %>%
  dplyr::filter(rowSums(.)!=0) %>%
  gather(sample, counts) %>%
  left_join(metadata, by = c("sample" = "sample")) %>%
  dplyr::mutate(date_of = as.factor(date_of)) %>%
  ggplot(aes(sample, log2(counts+1), fill = date_of)) +
  geom_boxplot() +
  coord_flip() +
  ggtitle("Counts per gene, protein_coding_genes")
```



Sample similarity analysis

Principal component analysis (PCA) - non zero genes


```
raw_counts <- assays(se)[["raw"]] %>%
  as_tibble() %>%
  dplyr::filter(rowSums(.)!=0) %>%
  as.matrix()

vst <- vst(raw_counts)

pca <- degPCA(vst, colData(se), condition = "date_of", name = "sample", data = T)[["plot"]]
pca_labels <- pca[["labels"]]
pca_data <- pca[["data"]] %>% as_tibble() %>%
  dplyr::select(sample, PC1, PC2, date_of)

pca_data$date_of <- as.factor(pca_data$date_of)

pca_data %>%
  ggplot(aes(x = PC1, y = PC2, color = date_of, label = sample)) +
  geom_point(size = 5) +
  geom_text_repel(size = 10) +
  xlab(pca_labels$x) +
  ylab(pca_labels$y) +
  ggtitle(paste0("PCA: ", nrow(vst), " genes")) +
  theme_bw(base_size = 20) +
  theme(axis.text = element_text(size = 20),
        axis.title = element_text(size = 20, face="bold"),
        legend.text = element_text(size = rel(1)))
```



Principal component analysis (PCA) - protein coding genes

```
raw_counts <- assays(se)[["raw"]] %>%
  as_tibble(rownames = "ensembl_gene_id") %>%
  dplyr::filter(ensembl_gene_id %in% protein_coding_genes$ensembl_gene_id) %>%
  column_to_rownames(var = "ensembl_gene_id") %>%
  dplyr::filter(rowSums(.) != 0) %>%
  as.matrix()

vst <- vst(raw_counts)

pca <- degPCA(vst, colData(se), condition = "date_of", name = "sample", data = T)[["plot"]]
pca_labels <- pca[["labels"]]
```


```

pca_data <- pca[["data"]] %>% as_tibble() %>%
  dplyr::select(sample, PC1, PC2, date_of)

pca_data$date_of <- as.factor(pca_data$date_of)

pca_data %>%
  ggplot(aes(x = PC1, y = PC2, color = date_of, label = sample)) +
  geom_point(size = 5) +
  geom_text_repel(size = 10) +
  xlab(pca_labels$x) +
  ylab(pca_labels$y) +
  ggtitle(paste0("PCA:", nrow(vst), " genes")) +
  theme_bw(base_size = 20) +
  theme(axis.text = element_text(size = 20),
        axis.title = element_text(size = 20, face="bold"),
        legend.text = element_text(size = rel(1)))

```

Principal component analysis (PCA) - protein coding genes - er

```
raw_counts <- assays(se)[["raw"]] %>%
  as_tibble(rownames = "ensembl_gene_id") %>%
  dplyr::filter(ensembl_gene_id %in% protein_coding_genes$ensembl_gene_id) %>%
  column_to_rownames(var = "ensembl_gene_id") %>%
  dplyr::filter(rowSums(.) != 0) %>%
  as.matrix()

vst <- vst(raw_counts)

pca <- degPCA(vst, colData(se), condition = "er", name = "sample", data = T)[["plot"]]
pca_labels <- pca[["labels"]]
```

```


pca_data <- pca[["data"]] %>% as_tibble() %>%
  dplyr::select(sample, PC1, PC2, er)

pca_data$er <- as.factor(pca_data$er)

pca_data %>%
  ggplot(aes(x = PC1, y = PC2, color = er, label = sample)) +
  geom_point(size = 5) +
  geom_text_repel(size = 10) +
  xlab(pca_labels$x) +
  ylab(pca_labels$y) +
  ggtitle(paste0("PCA:", nrow(vst), " genes")) +
  theme_bw(base_size = 20) +
  theme(axis.text = element_text(size = 20),
        axis.title = element_text(size = 20, face="bold"),
        legend.text = element_text(size = rel(1)))

```

PCA:19476 genes



PCA - protein coding genes - PRE/POST - Responce

PRE


```
se_pre <- se[se$treatment == "pre"]
raw_counts <- assays(se_pre)[["raw"]] %>%
  as_tibble(rownames = "ensembl_gene_id") %>%
  dplyr::filter(ensembl_gene_id %in% protein_coding_genes$ensembl_gene_id) %>%
  column_to_rownames(var = "ensembl_gene_id") %>%
  dplyr::filter(rowSums(.) != 0) %>%
  as.matrix()

vst <- vst(raw_counts)

pca <- degPCA(vst, colData(se_pre), condition = "response", name = "sample", data = T)[["plot"]]
pca_labels <- pca[["labels"]]
pca_data_pre <- pca[["data"]] %>% as_tibble() %>%
  dplyr::select(sample, PC1, PC2, response)


pca_data_pre$response <- as.factor(pca_data_pre$response)

pca_data_pre %>%
  ggplot(aes(x = PC1, y = PC2, color = response, label = sample)) +
  geom_point(size = 5) +
  geom_text_repel(size = 10) +
  xlab(pca_labels$x) +
  ylab(pca_labels$y) +
  ggtitle(paste0("PCA:", nrow(vst), " genes")) +
  theme_bw(base_size = 20) +
  theme(axis.text = element_text(size = 20),
        axis.title = element_text(size = 20, face="bold"),
        legend.text = element_text(size = rel(1)))
```



PRE - ROC

```
#pca_data$PC2
#pca_data$response
roc_pre <- roc(pca_data_pre$response, pca_data_pre$PC2, percent = FALSE)
ggroc(roc_pre,
      colour = "red", linetype = 2, size = 1, legacy.axes = TRUE)+
  theme_minimal() +
  scale_x_continuous(labels = scales::percent) +
  scale_y_continuous(labels = scales::percent) +
  xlab("100% - specificity : False Positive Percentage") +
  ylab("Sensitivity : True Postive Percentage") +
  ggtitle(paste0("PRE AUC : ", round(roc_pre$auc, 2)))
```



POST


```
se_post <- se[,se$treatment == "post"]
raw_counts <- assays(se_post)[["raw"]] %>%
  as_tibble(rownames = "ensembl_gene_id") %>%
  dplyr::filter(ensembl_gene_id %in% protein_coding_genes$ensembl_gene_id) %>%
  column_to_rownames(var = "ensembl_gene_id") %>%
  dplyr::filter(rowSums(.)!=0) %>%
  as.matrix()

vst <- vst(raw_counts)

pca_post <- degPCA(vst, colData(se_post), condition = "response", name = "sample", data = T)[["plot"]]
pca_labels <- pca_post[["labels"]]
pca_data_post <- pca_post[["data"]] %>% as_tibble() %>%
  dplyr::select(sample, PC1, PC2, response)

pca_data_post$response <- as.factor(pca_data_post$response)

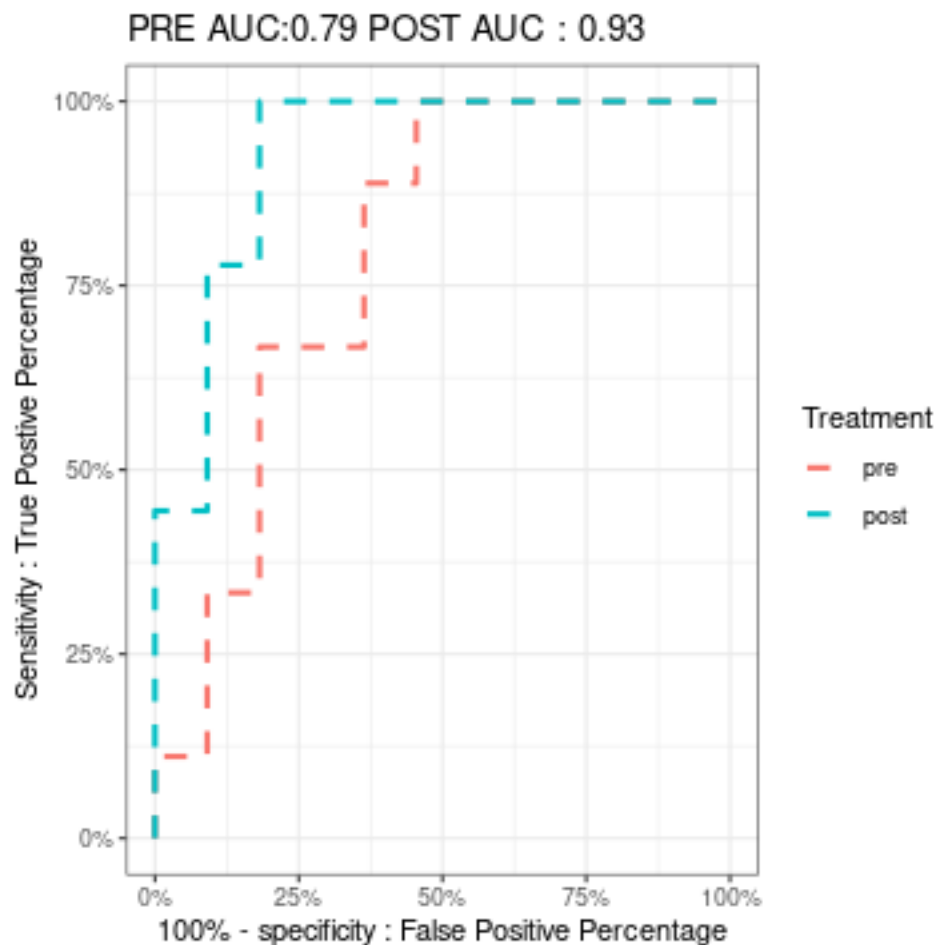
pca_data_post %>%
  ggplot(aes(x = PC1, y = PC2, color = response, label = sample)) +
  geom_point(size = 5) +
  geom_text_repel(size = 15) +
  xlab(pca_labels$x) +
  ylab(pca_labels$y) +
  ggtitle(paste0("PCA:", nrow(vst), " genes")) +
  theme_bw(base_size = 20) +
  theme(axis.text = element_text(size = 20),
        axis.title = element_text(size = 20, face="bold"),
        legend.text = element_text(size = rel(1)))
```



POST - ROC

```
#pca_data_post$PC2
#pca_data_post$response
# https://sachsmc.github.io/plotROC/
roc_post <- roc(pca_data_post$response, pca_data_post$PC2, percent = FALSE)

pROC::ggroc(list(pre = roc_pre,
  post = roc_post),
  legacy.axes = TRUE, size = 1, linetype = 2) +
  scale_x_continuous(labels = scales::percent) +
  scale_y_continuous(labels = scales::percent) +
  xlab("100% - specificity : False Positive Percentage") +
  ylab("Sensitivity : True Postive Percentage") +
  theme_bw() +
  ggtitle(paste0("PRE AUC:", round(roc_pre$auc, 2), " POST AUC : ", round(roc_post$auc, 2))) +
  guides(color = guide_legend(title="Treatment"))
```



R session

```
sessionInfo()
```

```
## R version 4.0.3 (2020-10-10)
## Platform: x86_64-redhat-linux-gnu (64-bit)
## Running under: Fedora 32 (Workstation Edition)
##
## Matrix products: default
## BLAS/LAPACK: /usr/lib64/libopenblas-r0.3.12.so
##
## locale:
##  [1] LC_CTYPE=en_CA.UTF-8      LC_NUMERIC=C
##  [3] LC_TIME=en_CA.UTF-8      LC_COLLATE=en_CA.UTF-8
##  [5] LC_MONETARY=en_CA.UTF-8  LC_MESSAGES=en_CA.UTF-8
##  [7] LC_PAPER=en_CA.UTF-8     LC_NAME=C
##  [9] LC_ADDRESS=C             LC_TELEPHONE=C
## [11] LC_MEASUREMENT=en_CA.UTF-8 LC_IDENTIFICATION=C
##
## attached base packages:
## [1] parallel stats4      stats      graphics  grDevices  utils      datasets
## [8] methods   base
##
## other attached packages:
##  [1] randomForest_4.6-14      pROC_1.17.0.1
##  [3] ggrepel_0.9.1            DEGREport_1.26.0
##  [5] DESeq2_1.30.1            SummarizedExperiment_1.20.0
##  [7] Biobase_2.50.0           MatrixGenerics_1.2.1
##  [9] matrixStats_0.58.0       GenomicRanges_1.42.0
## [11] GenomeInfoDb_1.26.7      IRanges_2.24.1
## [13] S4Vectors_0.28.1         BiocGenerics_0.36.1
## [15] knitr_1.30               forcats_0.5.1
## [17] stringr_1.4.0            dplyr_1.0.5
## [19] purrr_0.3.4              readr_1.4.0
## [21] tidyr_1.1.3              tibble_3.1.1
## [23] ggplot2_3.3.3            tidyverse_1.3.1
##
## loaded via a namespace (and not attached):
##  [1] colorspace_2.0-0         rjson_0.2.20
##  [3] ellipsis_0.3.1          circlize_0.4.12
##  [5] XVector_0.30.0          ggdendro_0.1.22
##  [7] GlobalOptions_0.1.2     fs_1.5.0
##  [9] clue_0.3-59             rstudioapi_0.13
## [11] farver_2.1.0            bit64_4.0.5
## [13] AnnotationDbi_1.52.0     fansi_0.4.2
## [15] lubridate_1.7.10        xml2_1.3.2
## [17] splines_4.0.3           logging_0.10-108
## [19] mnormt_2.0.2            cachem_1.0.4
## [21] geneplotter_1.68.0      jsonlite_1.7.2
## [23] Nozzle.R1_1.1-1         Cairo_1.5-12.2
## [25] broom_0.7.6             annotate_1.68.0
## [27] cluster_2.1.0           dbplyr_2.1.1
## [29] png_0.1-7              compiler_4.0.3
## [31] httr_1.4.2              backports_1.2.1
```

## [33] assertthat_0.2.1	Matrix_1.2-18
## [35] fastmap_1.1.0	limma_3.46.0
## [37] cli_2.5.0	lasso2_1.2-21.1
## [39] htmltools_0.5.1.1	tools_4.0.3
## [41] gtable_0.3.0	glue_1.4.2
## [43] GenomeInfoDbData_1.2.4	Rcpp_1.0.6
## [45] cellranger_1.1.0	vctrs_0.3.7
## [47] nlme_3.1-149	psych_2.1.3
## [49] xfun_0.19	rvest_1.0.0
## [51] lifecycle_1.0.0	XML_3.99-0.6
## [53] edgeR_3.32.1	MASS_7.3-53
## [55] zlibbioc_1.36.0	scales_1.1.1
## [57] hms_1.0.0	RColorBrewer_1.1-2
## [59] ComplexHeatmap_2.6.2	yaml_2.2.1
## [61] memoise_2.0.0	reshape_0.8.8
## [63] stringi_1.5.3	RSQLite_2.2.7
## [65] genefilter_1.72.1	BiocParallel_1.24.1
## [67] shape_1.4.5	rlang_0.4.10
## [69] pkgconfig_2.0.3	bitops_1.0-7
## [71] evaluate_0.14	lattice_0.20-41
## [73] labeling_0.4.2	cowplot_1.1.1
## [75] bit_4.0.4	tidyselect_1.1.0
## [77] plyr_1.8.6	magrittr_2.0.1
## [79] R6_2.5.0	generics_0.1.0
## [81] DelayedArray_0.16.3	DBI_1.1.1
## [83] pillar_1.6.0	haven_2.4.1
## [85] withr_2.4.2	survival_3.2-7
## [87] RCurl_1.98-1.3	modelr_0.1.8
## [89] crayon_1.4.1	utf8_1.2.1
## [91] tmvnsim_1.0-2	rmarkdown_2.5
## [93] GetoptLong_1.0.5	locfit_1.5-9.4
## [95] grid_4.0.3	readxl_1.3.1
## [97] blob_1.2.1	ConsensusClusterPlus_1.54.0
## [99] reprex_2.0.0	digest_0.6.27
## [101] xtable_1.8-4	munsell_0.5.0